Introduction

(Data Mining: Method and Application)

徐华

xuhua@tsinghua.edu.cn

1

Motivation - Background

- **The Explosive Growth of Data: from TeraBytes (TB) to PetaBytes (PB): B,KB,MB,GB,TB,PB(Big Data),EB,ZB,YB,DB,NB**
  - **Data collection and data availability**
  - **Ex. The changing flow of WebPages in China (2020-12)**

The Webpage Flow in China from CNNIC(2020-12)

2

Motivation - Background

- **About Big Data**
  - **> 1PB**
  - **“Big Data” is data whose scale, diversity, and complexity require new architecture, techniques, algorithms, and analytics to manage it and extract value and hidden knowledge from it…**
  - **Features: “5V”**
    - **Volume (规模大)**
    - **Variety （种类繁多）**
    - **Velocity （速度快）**
    - **Veracity （不确定性）**
    - **Value （价值）**

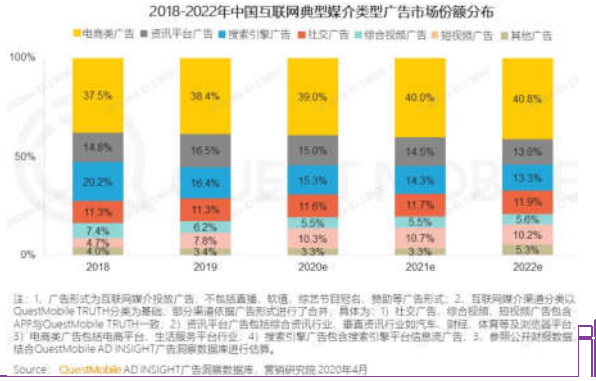

3

Motivation - Background

- **Features of Big Data**

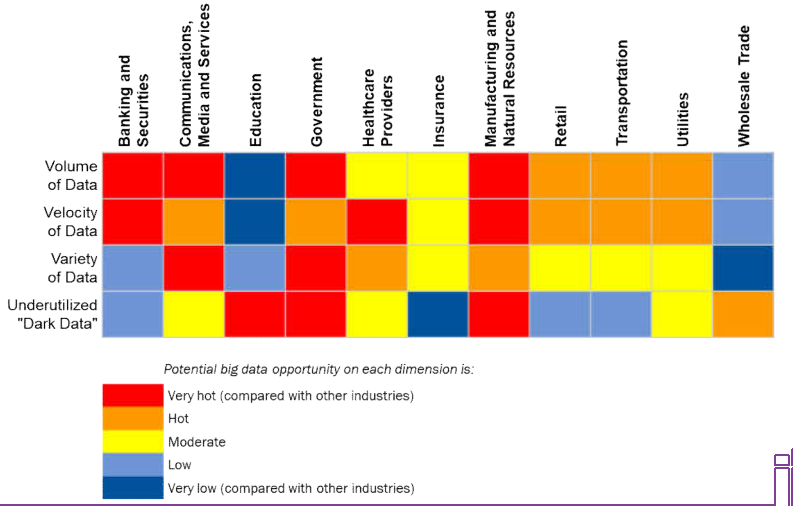

4

Motivation-Commercial Viewpoints

- **Commercial Viewpoints**
  - Data Sources: Web data, e-commerce, purchases at department/grocery stores, Bank/Credit Card, transactions
  - Computers have become cheaper and more powerful
  - Competitive Pressure is Strong
  - Provide better, customized services for an edge (e.g. in Customer Relationship Management)

5

Motivation – Scientific Viewpoints

- **Scientific Viewpoints**
  - Data collected and stored at enormous speeds (GB/hour)
    - remote sensors on a satellite
    - telescopes scanning the skies
    - microarrays generating gene expression data
    - scientific simulations generating terabytes of data
  - Traditional techniques infeasible for raw data
  - Data mining may help scientists
    - in classifying and segmenting data
    - in Hypothesis Formation

6

Motivation: Why Data mining?

- There is often information "hidden" in the data that is not readily evident
- Human analysts may take weeks to discover useful information
- Much of the data is never analyzed at all. " We are drowning in data, but starving for knowledge ! "
- "Necessity is the mother of invention" —Data mining—Automated analysis of massive data sets

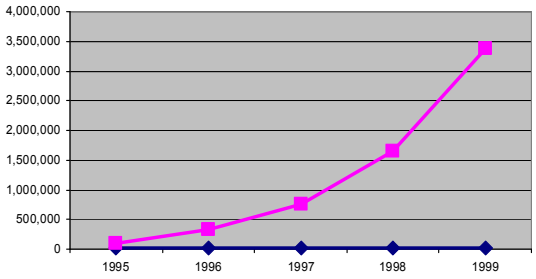

From: R. Grossman, C. Kamath, V. Kumar, "Data Mining for Scientific and Engineering Applications"

7

Evolution of Database Technology

- 1960s:
  - Data collection, database creation, IMS and network DBMS
- 1970s:
  - Relational data model, relational DBMS implementation
- 1980s:
  - RDBMS, advanced data models (extended-relational, OO, deductive, etc.)
  - Application-oriented DBMS (spatial, scientific, engineering, etc.)
- 1990s:
  - Data mining, data warehousing, multimedia databases, and Web databases
- 2000s
  - Stream data management and mining
  - Data mining and its applications
  - Web technology (XML, data integration) and global information systems

8

What is data mining?

- Data mining (knowledge discovery from data)
  - Extraction of interesting (non-trivial(非平凡的), implicit(隐含的), previously unknown(事先未知) and potentially useful(潜在用途)) patterns or knowledge from huge amount of data
  - Data mining: a misnomer?
- Alternative names
  - Knowledge Discovery (mining) in databases (KDD)
  - Knowledge Extraction
  - Data/pattern Analysis
  - Data Archeology (数据考古)
  - Data Dredging (数据捕捞/挖掘)
  - Information Harvesting
  - Business Intelligence
- Watch out: Is everything "data mining" ?
  - Simple search and query processing
  - (Deductive) expert systems

9

Knowledge Discovery Process

- Data mining—Core of Knowledge Discovery Process

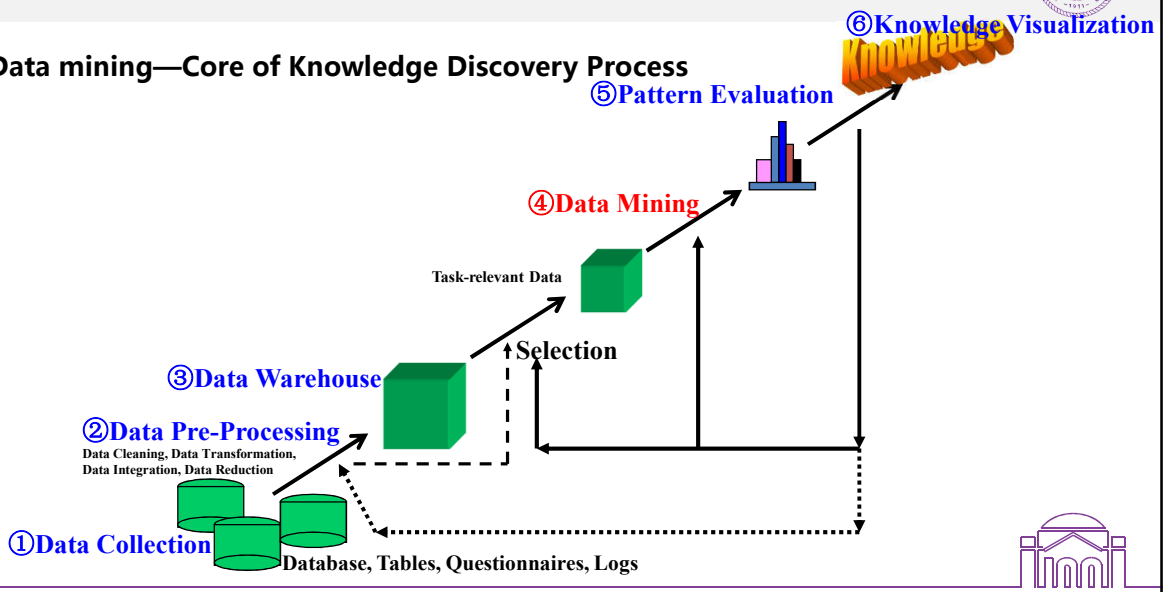

10

Data Mining v.s. KDD

- **Knowledge Discovery in Databases (KDD):** process of finding useful information and patterns in data.
- **Data Mining:** Use of algorithms to extract the information and patterns derived by the KDD process.

11

Data Mining and Business Intelligence

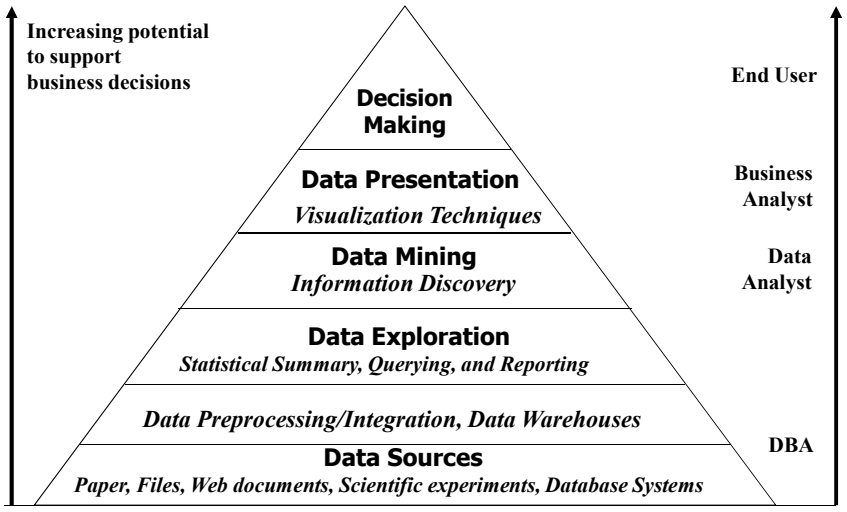

12

Data Mining: Confluence of Multiple Disciplines

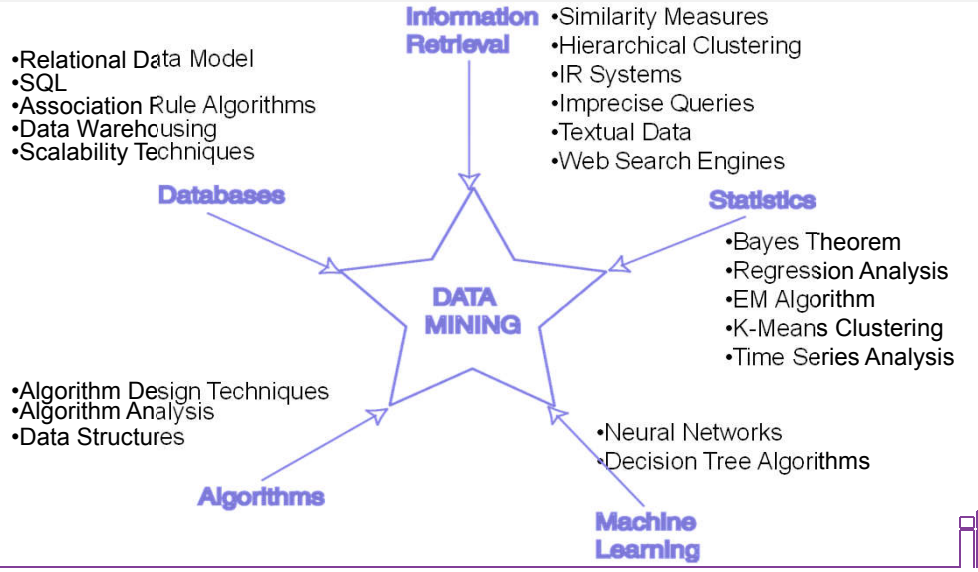

13

Why not traditional data analysis?

- **Tremendous amount of data**
  - Algorithms must be highly scalable to handle such as tera-bytes of data
- **High-dimensionality of data**
  - Micro-array may have tens of thousands of dimensions
- **High complexity of data**
  - Data streams and sensor data
  - Time-series data, temporal data, sequence data
  - Structure data, graphs, social networks and multi-linked data
  - Heterogeneous databases and legacy databases
  - Spatial, spatiotemporal, multimedia, text and Web data
  - Software programs, scientific simulations
- **New and sophisticated applications**

14

Multi-Dimensional View of Data Mining

- Data to be mined
  - Relational, data warehouse, transactional, stream, object-oriented/relational, active, spatial, time-series, text, multi-media, heterogeneous, legacy, WWW
- Knowledge to be mined
  - Characterization, discrimination, association, classification, clustering, trend/deviation, outlier analysis, etc.
  - Multiple/integrated functions and mining at multiple levels
- Techniques utilized
  - Database-oriented, data warehouse (OLAP), machine learning, statistics, visualization, etc.
- Applications adapted
  - Retail, telecommunication, banking, fraud analysis, bio-data mining, stock market analysis, text mining, Web mining, etc.

15

Data Mining: Classification Schemes

- General functionality
  - Descriptive data mining
  - Predictive data mining
- Different views lead to different classifications
  - Data view: Kinds of data to be mined
  - Knowledge view: Kinds of knowledge to be discovered
  - Method view: Kinds of techniques utilized
  - Application view: Kinds of applications adapted

16

Data Mining Functionalities(1)

- **Multidimensional concept description: Characterization and discrimination**
  - Generalize, summarize, and contrast data characteristics, e.g., Dry v.s. Wet regions
- **Frequent patterns, association, correlation vs. causality**
  - Diaper → Beer [0.5%, 75%] (Correlation or causality?)
- **Classification and prediction**
  - Construct models (functions) that describe and distinguish classes or concepts for future prediction
    - E.g., classify countries based on (climate), or classify cars based on (gas mileage)
  - Predict some unknown or missing numerical values

17

Data Mining Functionalities(2)

- **Cluster analysis**
  - Class label is unknown: Group data to form new classes, e.g., cluster houses to find distribution patterns
  - Maximizing intra-class similarity & minimizing interclass similarity
- **Outlier(离群点) analysis**
  - Outlier: Data object that does not comply with the general behavior of the data
  - Noise or exception? Useful in fraud detection, rare events analysis
- **Trend and evolution analysis**
  - Trend and deviation: e.g., regression analysis
  - Sequential pattern mining: e.g., digital camera -> large SD memory
  - Periodicity analysis
  - Similarity-based analysis
- **Other pattern-directed or statistical analysis**

18

About Top-10 DM Algorithms(1)

- **Classification**
  - #1. C4.5: Quinlan, J. R. C4.5: Programs for Machine Learning. Morgan Kaufmann., 1993.
  - #2. CART: L. Breiman, J. Friedman, R. Olshen, and C. Stone. Classification and Regression Trees. Wadsworth, 1984.
  - #3. K Nearest Neighbours (kNN): Hastie, T. and Tibshirani, R. 1996. Discriminant Adaptive Nearest Neighbor Classification. TPAMI. 18(6)
  - #4. Naive Bayes Hand, D.J., Yu, K., 2001. Idiot's Bayes: Not So Stupid After All? Internat. Statist. Rev. 69, 385-398.
- **Statistical Learning**
  - #5. SVM: Vapnik, V. N. 1995. The Nature of Statistical Learning Theory. Springer-Verlag.
  - #6. EM: McLachlan, G. and Peel, D. (2000). Finite Mixture Models. J. Wiley, New York. Association Analysis
  - #7. Apriori: Rakesh Agrawal and Ramakrishnan Srikant. Fast Algorithms for Mining Association Rules. In VLDB '94.
  - #8. FP-Tree: Han, J., Pei, J., and Yin, Y. 2000. Mining frequent patterns without candidate generation. In SIGMOD '00.

About Top-10 DM Algorithms(2)

- **Link Mining**
  - #9. PageRank: Brin, S. and Page, L. 1998. The anatomy of a large-scale hypertextual Web search engine. In WWW-7, 1998.
  - #10. HITS: Kleinberg, J. M. 1998. Authoritative sources in a hyperlinked environment. SODA, 1998.
- **Clustering**
  - #11. K-Means: MacQueen, J. B., Some methods for classification and analysis of multivariate observations, in Proc. 5th Berkeley Symp. Mathematical Statistics and Probability, 1967.
  - #12. BIRCH: Zhang, T., Ramakrishnan, R., and Livny, M. 1996. BIRCH: an efficient data clustering method for very large databases. In SIGMOD '96.
- **Bagging and Boosting**
  - #13. AdaBoost: Freund, Y. and Schapire, R. E. 1997. A decision-theoretic generalization of on-line learning and an application to boosting. J. Comput. Syst. Sci. 55, 1 (Aug. 1997), 119-139.

About Top-10 DM Algorithms(3)

- **Sequential Patterns**
  - #14. GSP: Srikant, R. and Agrawal, R. 1996. Mining Sequential Patterns: Generalizations and Performance Improvements. In Proceedings of the 5th International Conference on Extending Database Technology, 1996.
  - #15. PrefixSpan: J. Pei, J. Han, B. Mortazavi-Asl, H. Pinto, Q. Chen, U. Dayal and M-C. Hsu. PrefixSpan: Mining Sequential Patterns Efficiently by Prefix-Projected Pattern Growth. In ICDE '01.
- **Integrated Mining**
  - #16. CBA: Liu, B., Hsu, W. and Ma, Y. M. Integrating classification and association rule mining. KDD-98.
- **Rough Sets**
  - #17. Finding reduct: Zdzislaw Pawlak, Rough Sets: Theoretical Aspects of Reasoning about Data, Kluwer Academic Publishers, Norwell, MA, 1992
- **Graph Mining**
  - #18. gSpan: Yan, X. and Han, J. 2002. gSpan: Graph-Based Substructure Pattern Mining. In ICDM '02.

21

About Top-10 DM Algorithms(4)

- **Selected at ICDM2007**
  - #1: C4.5 (61 votes)
  - #2: K-Means (60 votes)
  - #3: SVM (58 votes)
  - #4: Apriori (52 votes)
  - #5: EM (48 votes)
  - #6: PageRank (46 votes)
  - #7: AdaBoost (45 votes)
  - #7: kNN (45 votes)
  - #7: Naive Bayes (45 votes)
  - #10: CART (34 votes)

22

Major Issues in Data Mining

- **Mining methodology**
  - Mining different kinds of knowledge from diverse data types, e.g., bio, stream, Web
  - Performance: efficiency, effectiveness, and scalability
  - Pattern evaluation: the interestingness problem
  - Incorporation of background knowledge
  - Handling noise and incomplete data
  - Parallel, distributed and incremental mining methods
  - Integration of the discovered knowledge with existing one: knowledge fusion
- **User interaction**
  - Data mining query languages and ad-hoc mining
  - Expression and visualization of data mining results
  - Interactive mining of knowledge at multiple levels of abstraction
- **Applications and social impacts**
  - Domain-specific data mining & invisible data mining
  - Protection of data security, integrity, and privacy

23

Textbook

Data Mining: Methods and Applications
(Second Edition)
by XU HUA
Tsinghua University Publishers, March 2022.
ISBN: 978-7-302-60144-9
ISBN: 978-7-302-36901-1

Data Mining: Methods and Applications—Application Examples
by XU HUA
Tsinghua University Publishers, Aug 2017.
ISBN: 978-7-302-47211-7

24

References(1)

Data Mining: Concepts and Techniques

by Jiawei Han & Micheline Kamber

Morgan Kaufmann Publishers, March 2006.

ISBN: 1-55860-901-6

数据挖掘教程

Margaret H.Dunham 著，郭崇慧，田凤占，靳晓明等译

清华大学出版社, 2005年5月.

ISBN: 7-302-10533-2

Data Mining: Introductory and Advanced Topics

by Margaret H. Dunham

Prentice Hall; Aug 2002

ISBN-10: 0130888923

25

References(2)

数据挖掘概念与技术（原书第3版）

韩家炜，堪博 著，范明，孟小峰 译

机械工业出版社，2007年3月

ISBN：9787111205388

Introduction to Data Mining

by Pang-Ning Tan, Michael Steinbach, Vipin Kumar

Addison Wesley; May 2005

ISBN-13: 978-0321321367

Web数据挖掘

Bing Liu 著，余勇，薛贵荣，韩定一 译

清华大学出版社，2009年4月

ISBN：978-7-302-19338-8

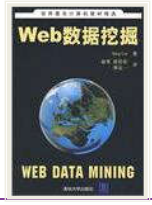

26

A Brief History of Data Mining Society

- 1989 IJCAI Workshop on Knowledge Discovery in Databases
  - Knowledge Discovery in Databases (G. Piatetsky-Shapiro and W. Frawley, 1991)
- 1991-1994 Workshops on Knowledge Discovery in Databases
  - Advances in Knowledge Discovery and Data Mining (U. Fayyad, G. Piatetsky-Shapiro, P. Smyth, and R. Uthurusamy, 1996)
- 1995-1998 International Conferences on Knowledge Discovery in Databases and Data Mining (KDD'95-98)
  - Journal of Data Mining and Knowledge Discovery (1997)
- ACM SIGKDD Conferences since 1998 and SIGKDD Explorations
- More conferences on data mining
  - PAKDD (1997), PKDD (1997), SIAM-Data Mining (2001), (IEEE) ICDM (2001), etc.
- ACM Transactions on KDD starting in 2007

27

Conferences

- KDD: ACM SIGKDD Int. Conf. on Knowledge Discovery in Databases and Data Mining
  - KDD08, KDD09 , KDD10 , KDD11 , KDD12, KDD13, KDD14, KDD15
- SDM: SIAM Data Mining Conf.
  - SDM08, SDM09, SDM10, SDM11, SDM12 , SDM13, SDM14, SDM15
- ICDM: IEEE Int. Conf. on Data Mining
  - ICDM08, ICDM09, ICDM10, ICDM11, ICDM12, ICDM13, ICDM14, ICDM15
- PKDD : Conf. on Principles and Practices of Knowledge Discovery and Data Mining
  - PKDD08, PKDD09, PKDD10, PKDD11, PKDD12, PKDD13, PKDD14, PKDD15
- PAKDD: Pacific-Asia Conf. on Knowledge Discovery and Data Mining
  - PAKDD08, PAKDD09, PAKDD10, PAKDD11, PAKDD12, PAKDD13,PAKDD14, PAKDD15

28

Journals

Data Mining and Knowledge Discovery

Geoffrey I. Webb

http://springer.lib.tsinghua.edu.cn/

清华IP可直接登录

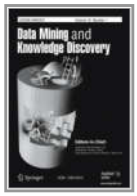

IEEE Trans. On Knowledge and Data Eng.

Xindong Wu

http://ieeexplore.ieee.org/xpl/RecentIssue.jsp?punumber=69

清华IP直接登录

SIGKDD Explorations

Osmar R. Zaiane

http://www.sigkdd.org/explorations/issue.php?issue=current

清华IP直接登录

29

Internet Resources(1)

- UCI数据集: http://kdd.ics.uci.edu/
- CMU数据集： http://lib.stat.cmu.edu/datasets/
  http://www.cs.cmu.edu/afs/cs.cmu.edu/project/theo-20/www/data/
- 时序数据集： http://www.stat.wisc.edu/~reinsel/bjr-data/
- 金融数据集： http://lisp.vse.cz/pkdd99/Challenge/chall.htm
- 癌症基因数据集： http://www.broadinstitute.org/cgi-bin/cancer/datasets.cgi
- 综合数据集：http://www.cs.nyu.edu/~roweis/data.html
- 数据集列表：http://www.kdnuggets.com/datasets/index.html
- 美国政府开放数据：http://data.gov
- 中国地方政府开放数据：北京 http://www.bjdata.gov.cn/
  上海 http://datashanghai.gov.cn

30

Internet Resources(2)

- UCI机器学习网站 http://archive.ics.uci.edu/ml/
- Weka官方网站 http://www.cs.waikato.ac.nz/ml/weka/
- DBMiner官方网站 http://ddm.cs.sfu.ca/
- SVM代码 http://www.csie.ntu.edu.tw/~cjlin/libsvm/
- 代码与数据集开源社区 https://github.com/
- 其它开源软件包： NB（朴素贝叶斯网络），NN（神经网络），DT（决策树）
- 相关软件：Matlab，StatSoft等商用软件；SQL Server 2008中也提供了相应的Data Analysis数据分析工具

31

Relative Courses

- Arizona
- Australian
- Bilkent
- CMU
- Central Connecticut
- Central Washington
- Cornel
- Depaul
- Georgia
- HKUST
- IIT, Indian
- McMaster
- Nanjing
- NUAA
- New York
- Pennsylvania
- Purdue
- RPI
- Rutgers
- Standford
- Alberta, Canada
- Wright State
- MIT
- Berkeley
- Helsinki, Finland
- Illinois
- Illinois at UC
- Massachusetts
- Minnesota
- Austin (1)(2)
- Toronto
- Washington
- Uppsala, Sweden
- VirginiaTech, USA

32

Summary

- Data mining: Discovering interesting patterns from large amounts of data
- A natural evolution of database technology, in great demand, with wide applications
- A KDD process includes data cleaning, data integration, data selection, transformation, data mining, pattern evaluation, and knowledge presentation
- Mining can be performed in a variety of information repositories
- Data mining functionalities: characterization, discrimination, association, classification, clustering, outlier and trend analysis, etc.
- Data mining systems and architectures
- Major issues in data mining

33

Recommended Reference Books

- S. Chakrabarti. Mining the Web: Statistical Analysis of Hypertex and Semi-Structured Data. Morgan Kaufmann, 2002
- R. O. Duda, P. E. Hart, and D. G. Stork, Pattern Classification, 2ed., Wiley-Interscience, 2000
- T. Dasu and T. Johnson. Exploratory Data Mining and Data Cleaning. John Wiley & Sons, 2003
- U. M. Fayyad, G. Piatetsky-Shapiro, P. Smyth, and R. Uthurusamy. Advances in Knowledge Discovery and Data Mining. AAAI/MIT Press, 1996
- U. Fayyad, G. Grinstein, and A. Wierse, Information Visualization in Data Mining and Knowledge Discovery, Morgan Kaufmann, 2001
- J. Han and M. Kamber. Data Mining: Concepts and Techniques. Morgan Kaufmann, 2nd ed., 2006
- D. J. Hand, H. Mannila, and P. Smyth, Principles of Data Mining, MIT Press, 2001
- T. Hastie, R. Tibshirani, and J. Friedman, The Elements of Statistical Learning: Data Mining, Inference, and Prediction, Springer-Verlag, 2001
- B. Liu, Web Data Mining, Springer 2006.
- T. M. Mitchell, Machine Learning, McGraw Hill, 1997
- G. Piatetsky-Shapiro and W. J. Frawley. Knowledge Discovery in Databases. AAAI/MIT Press, 1991
- P.-N. Tan, M. Steinbach and V. Kumar, Introduction to Data Mining, Wiley, 2005
- S. M. Weiss and N. Indurkhya, Predictive Data Mining, Morgan Kaufmann, 1998
- I. H. Witten and E. Frank, Data Mining: Practical Machine Learning Tools and Techniques with Java Implementations, Morgan Kaufmann, 2nd ed. 2005

34

概述部分结束！