

如何实现



- 常用的深度学习框架



61



如何实现



- 三个步骤



62



如何实现



● 手写体识别：定义网络、损失函数、优化器

```
# Define model
class NeuralNetwork(nn.Module):
    def __init__(self):
        super(NeuralNetwork, self).__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10)
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

model = NeuralNetwork().to(device)
print(model)
```

```
Using cuda device
NeuralNetwork(
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (linear_relu_stack): Sequential(
    (0): Linear(in_features=784, out_features=512, bias=True)
    (1): ReLU()
    (2): Linear(in_features=512, out_features=512, bias=True)
    (3): ReLU()
    (4): Linear(in_features=512, out_features=10, bias=True)
  )
)
```

```
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=1e-3)
```

如何实现



● 手写体识别：数据准备

```
# Download training data from open datasets.
training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor(),
)

# Download test data from open datasets.
test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor(),
)
```

```
batch_size = 64

# Create data loaders.
train_dataloader = DataLoader(training_data, batch_size=batch_size)
test_dataloader = DataLoader(test_data, batch_size=batch_size)

for X, y in test_dataloader:
    print(f"Shape of X [N, C, H, W]: {X.shape}")
    print(f"Shape of y: {y.shape} {y.dtype}")
    break
```

Out:

```
Shape of X [N, C, H, W]: torch.Size([64, 1, 28, 28])
Shape of y: torch.Size([64]) torch.int64
```



如何实现

● 手写体识别：模型训练和优化

```
def train(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
    model.train()
    for batch, (X, y) in enumerate(dataloader):
        X, y = X.to(device), y.to(device)

        # Compute prediction error
        pred = model(X)
        loss = loss_fn(pred, y)

        # Backpropagation
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    if batch % 100 == 0:
        loss, current = loss.item(), batch * len(X)
        print(f"loss: {loss:>7f}    [{current:>5d}/{size:>5d}]")
```

65

Epoch 1

```
loss: 2.290059 [ 0/60000]
loss: 2.278689 [ 6400/60000]
loss: 2.260067 [12800/60000]
loss: 2.260366 [19200/60000]
loss: 2.251141 [25600/60000]
loss: 2.207952 [32000/60000]
loss: 2.221357 [38400/60000]
loss: 2.185180 [44800/60000]
loss: 2.179188 [51200/60000]
loss: 2.156274 [57600/60000]
Test Error:
Accuracy: 50.4%, Avg loss: 2.14366
```

Epoch 2

```
loss: 2.147471 [ 0/60000]
loss: 2.136722 [ 6400/60000]
loss: 2.077863 [12800/60000]
loss: 2.101987 [19200/60000]
loss: 2.059406 [25600/60000]
loss: 1.982178 [32000/60000]
loss: 2.019614 [38400/60000]
loss: 1.936746 [44800/60000]
loss: 1.944287 [51200/60000]
loss: 1.871792 [57600/60000]
Test Error:
Accuracy: 61.9%, Avg loss: 1.867935
```

Epoch 5

```
loss: 1.300350 [ 0/60000]
loss: 1.288346 [ 6400/60000]
loss: 1.112341 [12800/60000]
loss: 1.230672 [19200/60000]
loss: 1.108364 [25600/60000]
loss: 1.116301 [32000/60000]
loss: 1.159760 [38400/60000]
loss: 1.088508 [44800/60000]
loss: 1.133090 [51200/60000]
loss: 1.046162 [57600/60000]
Test Error:
Accuracy: 65.1%, Avg loss: 1.066427
```




如何实现

● 三个步骤



Deep Learning is so simple



66

CDC.TENCENT.COM



如何实现



- 实现很简单！



67



谢谢！

68

